

INSE 6421 System Integration and Testing
TV Script Generation using RNN

Report By
Binit Kumar (40172005)
Meghaksh Brahmbhatt (40166917)

April 19, 2022

Abstract

Artificial intelligence has totally added a lot of new things in the domain of possibilities that were not possible a while ago. It has slowly evolved into an industry that encompasses everything from banking and medical to music, gaming, and a variety of other fields.

A neural network is a type of computational network. One of the most advanced systems is one that is based on the human brain. In this case, We implemented a type of neural network called recurrent neural network that works particularly well for sequential data and suitable for projects like TV Script generation. Further we have also discussed the background of RNN, issues with simple RNN, their solutions and comparisons between different RNN architectures and approaches.

Tv script generation is a text creation project that is based on LSTM and GRU architecture. The goal of our system is to create a new fake script based on a series of script from a popular Tv series called Seinfeld.[\[8\]](#) We have analyzed various neural network architectures and libraries such as pytorch to create the best model to implement LSTM and eventually create a functional prototype. All the codes are available in Github.[\[6\]](#) The system's output is evaluated using parameters such as grammar correctness, event linkage, and uniqueness.

Contents

List of Figures	ii
1 Introduction	1
1.1 Project Details	1
1.2 Problem Statement	1
1.3 Methodology	1
1.4 Assumptions and Limitations	1
1.5 Input and Expected Output Script	2
2 Background	3
2.1 Why using RNN?	3
2.2 Visualizing Character-wise RNN	3
2.3 Problem with RNN Architecture	5
2.3.1 Vanishing and Exploding Gradients	6
3 Overcoming the issues with RNN	7
3.1 RNN Cell	7
3.2 LSTM Cell	7
3.2.1 Forget Gate	8
3.2.2 Update Gate	8
3.2.3 Cell state to hidden output	9
3.2.4 Putting it all together	9
3.3 How gradient problem is fixed?	10
4 Implementation Details	11
4.1 Overall Approach	11
4.2 Framework/Library Used	12
4.3 Pre-processing of Input Data	12
4.3.1 Tokenize Punctuations	13
4.4 Creating batches of data	13
4.5 How data-loader looks like?	14
4.6 Neural Network Architecture	15
4.7 Forward Propagation	16
4.8 Backpropagation	17
4.9 Training Loop	17
4.10 Hyperparameters	17

4.11	Training Results	20
4.12	Final output	20
5	Conclusions	22
5.1	Comparision of Different Architectures	22
5.2	Final Comments	22
	Bibliography	23

List of Figures

1.1	Sample Input and Expected output script [6]	2
2.1	Predicting next character in the word "steep" [5]	4
2.2	Understanding the hidden layer [5]	4
2.3	Unrolling the network [5]	5
2.4	One hot encoding [5]	5
2.5	Understanding gradient problem [5]	6
2.6	Vanishing and Exploding Gradients [5]	6
3.1	Basic RNN Cell [5]	7
3.2	Forget Gate [5]	8
3.3	Update Gate [5]	8
3.4	Cell state to hidden output [5]	9
3.5	Simple LSTM Cell [5]	9
3.6	Fixing Gradient Problem [5]	10
4.1	Various steps in implementation [2]	11
4.2	Pre-processing input data [6]	12
4.3	Tokenization [6]	13
4.4	Creating batches of data [6]	14
4.5	data-loader as a tensor [6]	14
4.6	Initializing RNN [6]	15
4.7	Initializing hidden state	15
4.8	Forward Propagation [6]	16
4.9	Implementing RNN using GRU [6]	17
4.10	Forward and Back Propagation in Training loop [6]	18
4.11	Training Loop [6]	19
4.12	Adam Optimizer and Cross Entropy [6]	19
4.13	Training progress result [6]	20
4.14	Generated Script using RNN [6]	21

1 Introduction

1.1 Project Details

A project to generate a fake Tv-script using a RNN model. We are using a Tv script from a popular Tv show named Seinfeld.[8] The neural network takes the following mentioned tv-script, learns the text sequences from it to generate a new fake tv-script. The input dataset is taken from Kaggle, named Seinfeld Chronicles. [7]

1.2 Problem Statement

The problem statement can be described as a unique requirement to save the sequential and semantic information of text. For e.g. any word that comes after a given word or before a given word defines its meaning in a unique context in the sequence. [4]Also, the idea is to figure out which neural network architectures is best for this project and creating a working prototype to showcase this.

1.3 Methodology

The architecture used is LSTM i.e. Long Short Term Memory is a type of RNN architecture. We will discuss further in the document about why other architecture are not suitable and what are the alternative architectures for this and comparison between them.[1]

1.4 Assumptions and Limitations

- We are assuming that the input tv-script will mostly be grammatically correct.
- The output generated script will be structured grammatically and will not be necessarily making sense all the time.
- This project is only done to show the implementation using LSTM. It is limited by the amount of data and training to show bare minimum working code.
- This is not a generic model to work for all types of TV Script but it is well suited for the given task only.

1.5 Input and Expected Output Script

Input Sample	Output Sample
george: so, i was in the contrary. george: so, i guess i was a woman, and the defendants was a good boy. elaine: i don't know where it is. jerry: so i was thinking about this one? hoyt: i thought i was a little adjustment. jerry: what is that? hoyt: yes, yes. yes. i got a pee on this. you were in the middle of a plane, and i have a little adjustment. george: what happened to him?	george: so you want to see how you could do that? hoyt: i don't want to see you in a hotel. elaine: you want to get the car on the street and a wheelchair? hoyt: what do you think? george: yeah, i guess i was wondering about it. jerry: what are you doing with this girl? elaine: no, i got to tell him. george: i can't believe this is the most exciting thing.

Figure 1.1: Sample Input and Expected output script [6]

2 Background

2.1 Why using RNN?

Our script generation takes a word and determine the next word in a sentence. This requires keeping a sequence or order of words in the neural network. However, In Feed-Forward Network, there is no sense of order in the input. So, the question is how to build the concept of memory in the neural network so that it can learning for the sequence. Recurrence relations need to apply at each time step t , such that model learns the concepts of memory by updating the hidden state.

$$h_t = f_w(h_{t-1}, x_t)$$

output vector can be calculated as,

$$y_t = g_w(h_t)$$

2.2 Visualizing Character-wise RNN

Our goal here is to predict the next character in word “steep” [5]

- When we pass “s”, desired output is “t”.
- When we pass “t”, desired output is “e”.
- When we pass “e”, desired output could be “e” or “p”.

The network doesn’t have enough information to determine which character to predict! To solve this problem, we need to include information about the sequence of characters.

We can solve this problem by routing the hidden layer output from the previous step back into the hidden layer.

- The box in the diagram means the value from the previous sequence, or time step.
- Now the network sees an “e”, it knows it saw an “s” and a “t” before, so the next character should be another “e”

The equation for the hidden layer can be given as,

$$h_t = f(W_{hh}h_{t-1} + W_{xh}x_t)$$

This architecture is known as Recurrent Neural Network or RNN




Figure 2.1: Predicting next character in the word "steep" [5]




Figure 2.2: Understanding the hidden layer [5]

- Now the total input in the hidden layer is the sum of the layered combinations from the input layer and previous hidden layer

We can view our recurrent network as one big graph by unrolling it.

- Now, we have a feed-forward network for each character but connected through the hidden layers.
- Each hidden nodes receives inputs from input node and hidden node from the previous step

Let's visualize by adding some numbers here.

- Here, we're one hot encoding the input characters.
- $1000 = \text{"s"} , 0100 = \text{"t"} , 0010 = \text{"e"} , 0010 = \text{"e"}$
- There are three units in the hidden layer and the output layer is showing the logits




Figure 2.3: Unrolling the network [5]

- We pass the logits into Softmax function to get prediction and to train with a cross entropy loss.

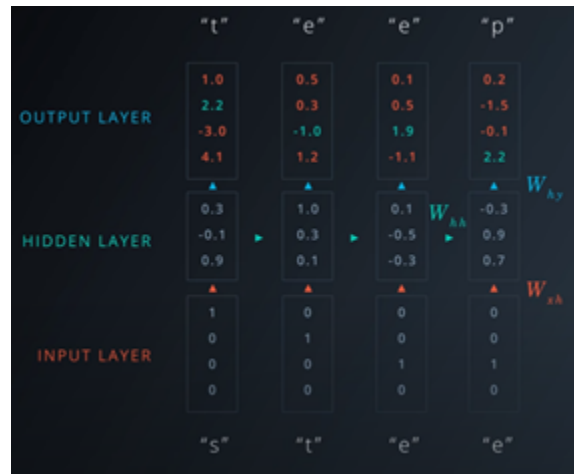


Figure 2.4: One hot encoding [5]

2.3 Problem with RNN Architecture

We can include information from a sequence of data using a recurrent connection on the hidden layer. This connection goes through these weights, W_{hh} . After enrolling the network, we say the hidden layer at step t is a function of the previous hidden state multiplied by those weights. The output of that layer is again multiplied by W_{hh} . For every step we have in the network, we are multiplying by the weights again and again. And when we do backpropagation, that's even more multiplication. These leads to problem where gradients going through network either get really small and vanish or get really large and explode.




Figure 2.5: Understanding gradient problem [5]

2.3.1 Vanishing and Exploding Gradients

If we multiply by some number a bunch of times, we will get two results except a couple of special cases. If that number is less than 1, we will end up at 0. If it greater than 1, we will head towards infinity. This happens to gradient in normal RNN, where they either vanish or explode. Resulting in making it difficult for RNNs to learn long range interactions.




Figure 2.6: Vanishing and Exploding Gradients [5]

3 Overcoming the issues with RNN

3.1 RNN Cell

We can think of RNNs as a bunch of cells with inputs and outputs. Inside the cell, we have network layers, such as the sigmoid layer labelled with a sigma here. To solve the problem of the vanishing gradients, we can use more complicated cells called long short-term memory or LSTM.

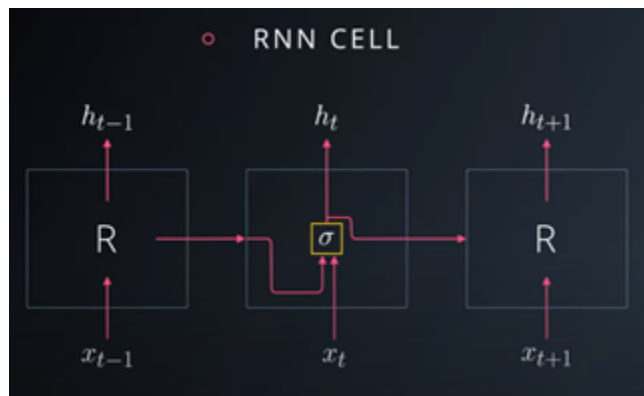


Figure 3.1: Basic RNN Cell [5]

3.2 LSTM Cell

For better understanding, LSTM Cell can be broken down as below: The key addition here is the cell state labelled C. In this cell, there are four network layers shown as yellow boxes. Each of them with their own weights. The layers labelled with sigma are sigmoid and tanh is the hyperbolic tangent function. Tanh is similar to a sigmoid in that it squashes input, but the output is between -1 to 1 instead of 0 and 1. The red circles are point-wise or element-wise operations i.e. they operate on matrices element by element. The main improvement here is through the cell state. The cell state goes through LSTM cell with little interaction allowing information to flow easily through the cells. The cell state is modified only through element-wise operation which functions as gates. And the hidden state is now calculated through cell state, then passed on. [4]

3.2.1 Forget Gate

The first gate is the forget gate. The values coming out of sigmoid layer are between 0 and 1. Then they are multiplied element-wise with the cell state. So the values from this layer close to 0 will shut off certain elements in the cell state. Effectively, forgetting that information going forward. Conversely, values close to 1 will allow information to pass through unchanged. This is helpful, because the network can learn to forget information that causes incorrect predictions. On the other hand, long range information that are helpful is allowed to flow through freely.

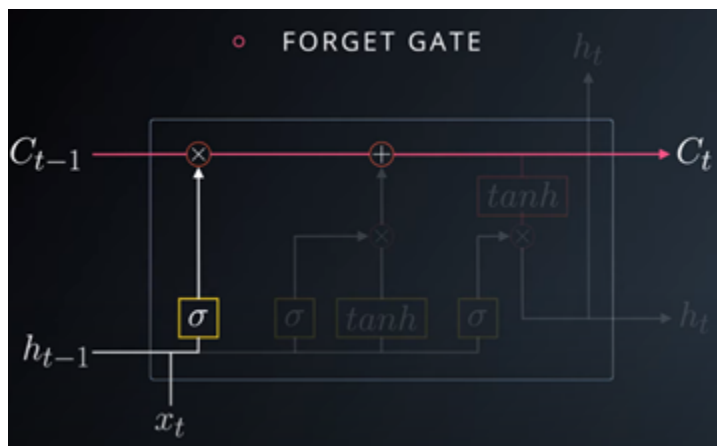


Figure 3.2: Forget Gate [5]

3.2.2 Update Gate

The next gate updates the cell state from the input and previous hidden state. The \tanh layer output is added to the cell state and again gated by a sigmoid layer. In this way, the cell state can be updated in the step and passed along to the next cell

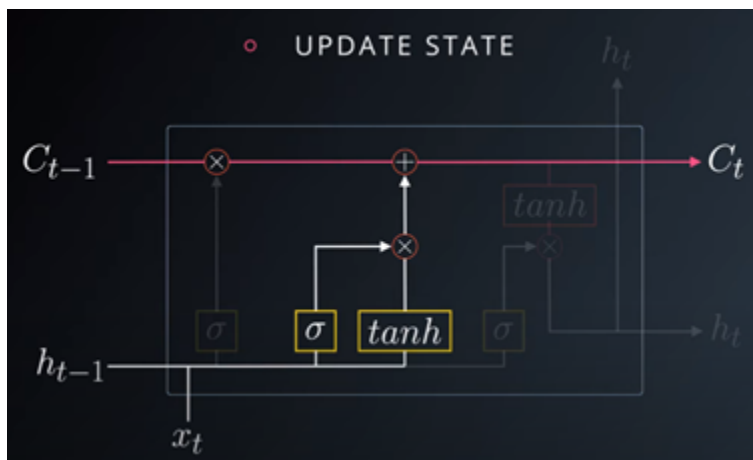


Figure 3.3: Update Gate [5]

3.2.3 Cell state to hidden output

Here, the cell state is used to produce the hidden state which is sent to the next hidden cell as well as to higher layers. It's the arrow pointing up here The cell state is passed through another tanh then gated again with another sigmoid layer. All these sigmoid gates let the network learn which information to keep and which information to get rid of.




Figure 3.4: Cell state to hidden output [5]

3.2.4 Putting it all together

Putting all this together, the LSTM cell consists of a cell state with a bunch of gates used to update it, and leak it out to the hidden state. This is just the basic LSTM. There are multiple variations and lot of ongoing experimentation into improving these. They are also stacked into deeper layer. We just send the output from one cell to the input of another

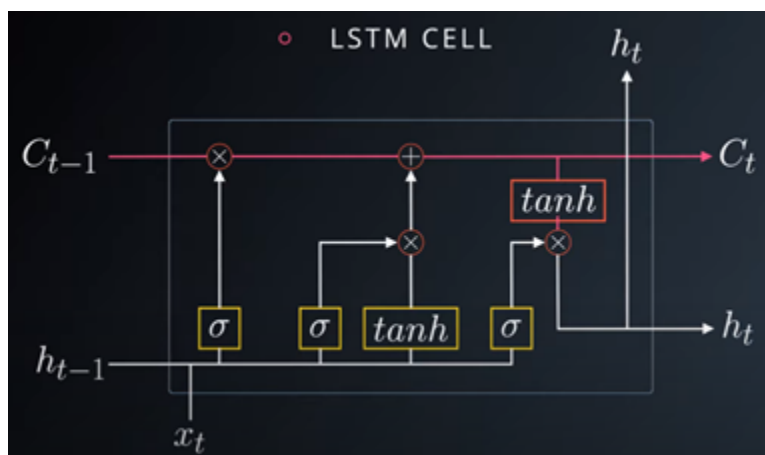


Figure 3.5: Simple LSTM Cell [5]

3.3 How gradient problem is fixed?

Since the cell state is allowed to flow through the hidden layers with only this linear sum operation. Gradient can easily move through the network without being diminished. We can also get gradients added into the network through the LSTM cells but they are just added to the gradients flowing through. LSTM are basic unit of RNN in many applications.



Figure 3.6: Fixing Gradient Problem [5]

4 Implementation Details

4.1 Overall Approach

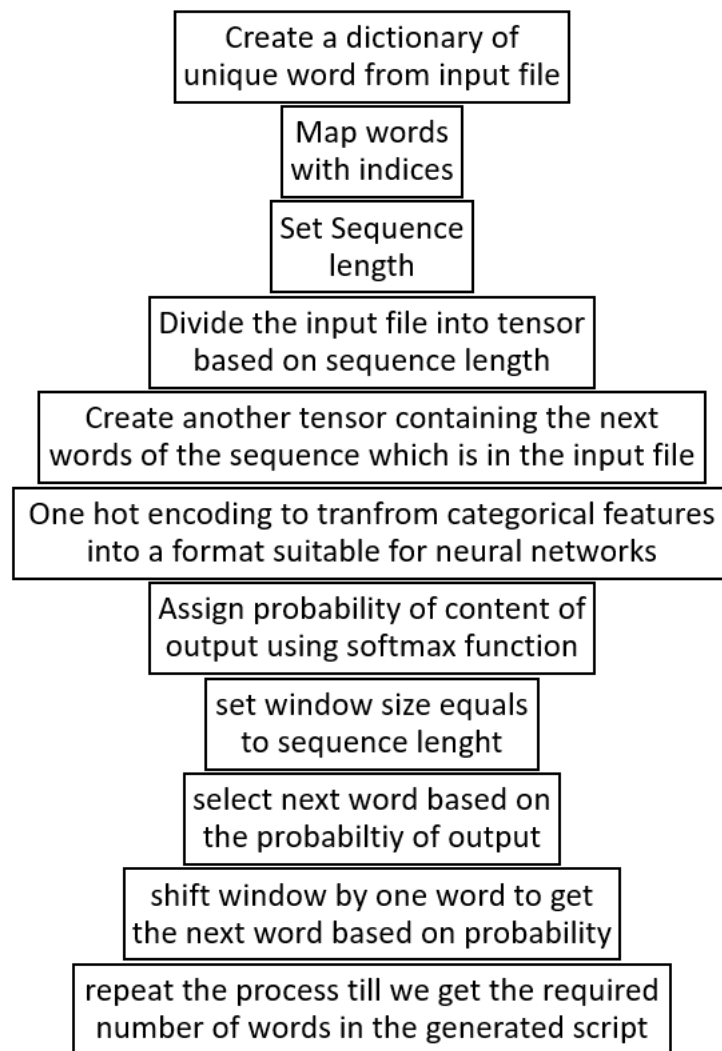


Figure 4.1: Various steps in implementation [2]

4.2 Framework/Library Used

- Pytorch – An open-source machine learning framework that accelerates the path from research prototyping to production deployment.[11]
- Numpy – A library that offers comprehensive mathematical functions, random number generators, linear algebra routines, Fourier transforms and more.[12]
- Pickle – A library to convert Python object hierarchy into byte stream called “pickling”. “unpickling” is the inverse operation where a byte stream is converted back into an object hierarchy.[13]

4.3 Pre-processing of Input Data

Pre-processing is a step before feeding the input to the neural network where the input data is pre-processed to make it suitable for training. This includes,

- Changing the entire data set into lowercase
- Splitting the sentences to get all the words
- Creating lookup table to generate word embeddings i.e. transforms the word to integer ids
- vocab-to-int: dictionary to go from a word to id
- int-to-vocab: dictionary to go from the id to word

```
def create_lookup_tables(text):  
    word_counts = Counter(text)  
    # sort words from most to least frequent in occurrence  
    sorted_vocab = sorted(word_counts, key = word_counts.get, reverse=True)  
    # create int_to_vocab dictionaries  
    int_to_vocab = {ii: word for ii, word in enumerate(sorted_vocab)}  
    vocab_to_int = {word: ii for ii, word in int_to_vocab.items()}  
  
    # return tuple  
    return (vocab_to_int, int_to_vocab)
```

Figure 4.2: Pre-processing input data [6]

The above function takes the tv script as input and return two dictionaries as tuple containing word to id and id to word entries.

4.3.1 Tokenize Punctuations

Punctuations like periods and exclamation marks can create multiple ids for the same word. For e.g. bye, bye! This dictionary will be used to tokenize the symbols and add the delimiter (space) around it. This separates each symbol as its own word, making it easier for the neural network to predict the next word.

```
def token_lookup():
    """
    Generate a dict to turn punctuation :
    :return: Tokenized dictionary where :
    """
    punct_dict = {'.': '<PERIOD>',
                  ',': '<COMMA>',
                  '"': '<QUOTATION_MARK>',
                  ';': '<SEMICOLON>',
                  '!': '<EXCLAMATION_MARK>',
                  '?': '<QUESTION_MARK>',
                  '(': '<LEFT_BRAC>',
                  ')': '<RIGHT_BRAC>',
                  '-': '<DASH>',
                  '\n': '<NEW_LINE>'
                  }
    return punct_dict
```

Figure 4.3: Tokenization [6]

4.4 Creating batches of data

batch_data function to batch words data into chunks of size batch_size using Pytorch's DataLoader classes. DataLoader class will help to create feature_tensors and target_tensors of correct size and content of given sequence_length . E.g. words = [1, 2, 3, 4, 5, 6, 7]

- sequence_length = 4
- First feature_tensor would be: [1, 2, 3, 4]
- The corresponding target_tensor would be: 5
- Second feature_tensor would be: [2, 3, 4, 5]
- And the second target_tensor would be: 6

The sample_x should be of size (batch_size, sequence_length) or (10, 5) in this case and sample_y should just have one dimension: batch_size (10). We can also notice that the target sample_y are the next value in the ordered test_text data.

```

from torch.utils.data import TensorDataset, DataLoader

def batch_data(words, sequence_length, batch_size):
    """
    Batch the neural network data using DataLoader
    :param words: The word ids of the TV scripts
    :param sequence_length: The sequence length of each batch
    :param batch_size: The size of each batch; the number of sequences in a batch
    :return: DataLoader with batched data
    """
    # DONE: Implement function
    feature_tensors, target_tensors = [], []
    for idx in range(0, len(words)-sequence_length):
        feature_tensors.append( words[idx: idx+sequence_length] )
        target_tensors.append( words[idx+sequence_length] )

    feature_tensors = torch.tensor(feature_tensors)
    target_tensors = torch.tensor(target_tensors)

    data = TensorDataset(feature_tensors, target_tensors)
    data_loader = DataLoader(data, batch_size=batch_size, shuffle=True)

    # return a dataloader
    return data_loader

```

Figure 4.4: Creating batches of data [6]

4.5 How data-loader looks like?

Sample batch of inputs `sample_x` and targets `sample_y` from the data loader. We are shuffling the data in the data loader to get random batches.

```

torch.Size([10, 5])
tensor([[ 28,  29,  30,  31,  32],
        [ 21,  22,  23,  24,  25],
        [ 17,  18,  19,  20,  21],
        [ 34,  35,  36,  37,  38],
        [ 11,  12,  13,  14,  15],
        [ 23,  24,  25,  26,  27],
        [  6,   7,   8,   9,  10],
        [ 38,  39,  40,  41,  42],
        [ 25,  26,  27,  28,  29],
        [  7,   8,   9,  10,  11]])

torch.Size([10])
tensor([ 33,  26,  22,  39,  16,  28,  11,  43,  30,  12])

```

Figure 4.5: data-loader as a tensor [6]

4.6 Neural Network Architecture

The below mentioned class RNN inherit the `nn.Module` class from Pytorch library. `__init__` method is the constructor of the class to initialize the Pytorch RNN module with parameters like `vocab_size`, `output_size`, `embedding_dim`, `hidden_dim` and `dropout`.

```
class RNN(nn.Module):  
  
    def __init__(self, vocab_size, output_size, embedding_dim, hidden_dim, n_layers, dropout=0.5):  
        """  
        Initialize the PyTorch RNN Module  
        :param vocab_size: The number of input dimensions of the neural network (the size of the vocabulary)  
        :param output_size: The number of output dimensions of the neural network  
        :param embedding_dim: The size of embeddings, should you choose to use them  
        :param hidden_dim: The size of the hidden layer outputs  
        :param dropout: dropout to add in between LSTM/GRU layers  
        """  
        super(RNN, self).__init__()  
        # TODO: Implement function  
  
        # set class variables  
        self.output_size = output_size  
        self.n_layers = n_layers  
        self.hidden_dim = hidden_dim  
  
        # define model layers  
        self.embedding = nn.Embedding(vocab_size, embedding_dim)  
        self.lstm = nn.LSTM(embedding_dim, hidden_dim, n_layers, dropout=dropout, batch_first=True)  
        self.dropout = nn.Dropout(dropout)  
        self.fc = nn.Linear(hidden_dim, output_size)
```

Figure 4.6: Initializing RNN [6]

```
def init_hidden(self, batch_size):  
    """  
    Initialize the hidden state of an LSTM/GRU  
    :param batch_size: The batch_size of the hidden state  
    :return: hidden state of dims (n_layers, batch_size, hidden_dim)  
    """  
    weight = next(self.parameters()).data  
  
    # initialize hidden state with zero weights, and move to GPU if available  
    if train_on_gpu:  
        hidden = (weight.new(self.n_layers, batch_size, self.hidden_dim).zero_().cuda(),  
                  weight.new(self.n_layers, batch_size, self.hidden_dim).zero_().cuda())  
    else:  
        hidden = (weight.new(self.n_layers, batch_size, self.hidden_dim).zero_(),  
                  weight.new(self.n_layers, batch_size, self.hidden_dim).zero_())  
  
    return hidden
```

Figure 4.7: Initializing hidden state

4.7 Forward Propagation

```
def forward(self, nn_input, hidden):
    """
    Forward propagation of the neural network
    :param nn_input: The input to the neural network
    :param hidden: The hidden state
    :return: Two Tensors, the output of the neural network and the latest hidden state
    """
    # get the batch size
    batch_size = nn_input.size(0)

    # Embedding and LSTM layer
    embeds = self.embedding(nn_input)
    lstm_out, hidden = self.lstm(embeds, hidden)

    # stack up LSTM output to pass through the last FC output layer
    lstm_out = lstm_out.contiguous().view(-1, self.hidden_dim)

    # dropout layer after LSTM
    output = self.dropout(lstm_out)

    # last FC layer
    output = self.fc(output)

    # reshape into (batch_size, seq_length, output_size)
    output = output.view(batch_size, -1, self.output_size)

    # get last batch
    out = output[:, -1]

    # return one batch of output word scores and the hidden state
    return out, hidden
```

Figure 4.8: Forward Propagation [6]

Forward method in the RNN class defines the forward propagation of the neural network. Shown in figure [4.8] and [4.7]

Alternatively, we can also use GRU (i.e. Gated Recurrent Units) as well in place of LSTM. GRU is related to LSTM as both are utilizing different way if gating information to prevent vanishing gradient problem. Shown in figure [4.9]

- The GRU controls the flow of information like the LSTM unit, but without having to use a memory unit. It just exposes the full hidden content without any control.
- GRU is relatively new and performance is similar to LSTM but computationally more efficient because of less complex structure

```

class RNN(nn.Module):

    def __init__(self, vocab_size, output_size, embedding_dim, hidden_dim, n_layers, dropout=0.5):

        super(RNN, self).__init__()

        # set class variables
        self.output_size = output_size
        self.n_layers = n_layers
        self.hidden_dim = hidden_dim

        # define model layers
        self.embedding = nn.Embedding(vocab_size, embedding_dim)
        self.gru = nn.GRU(embedding_dim, hidden_dim, n_layers, batch_first=True, dropout=dropout)
        self.fc = nn.Linear(hidden_dim, output_size)

    def forward(self, nn_input, hidden):

        batch_size = nn_input.size(0)
        embeds = self.embedding(nn_input)
        gru_out, hidden = self.gru(embeds, hidden)

        # stack up gru outputs
        gru_out = gru_out.contiguous().view(-1, self.hidden_dim)

        # fully connected layer
        output = self.fc(gru_out)

        # reshape into (batch_size, seq_length, output_size)
        output = output.view(batch_size, -1, self.output_size)
        # get last batch
        output = output[:, -1]

```

Figure 4.9: Implementing RNN using GRU [6]

4.8 Backpropagation

Below function will be called iteratively in the training loop. This applies forward and back propagation. Shown in figure [4.10]

4.9 Training Loop

This function train the network over all the batches for the number of epochs given. The model progress is printed every number of batches. Shown in figure [4.11]

4.10 Hyperparameters

The values of the below parameters can be found here.[6]

- Sequence Length: to set the length of the sequence
- Batch Size : to set the size of batches
- Number of epochs : to set the number of epochs to train for

```

def forward_back_prop(rnn, optimizer, criterion, inp, target, hidden):
    """
    Forward and backward propagation on the neural network
    :param decoder: The PyTorch Module that holds the neural network
    :param decoder_optimizer: The PyTorch optimizer for the neural network
    :param criterion: The PyTorch loss function
    :param inp: A batch of input to the neural network
    :param target: The target output for the batch of input
    :return: The loss and the latest hidden state Tensor
    """

    # move data to GPU, if available
    if train_on_gpu:
        inp, target = inp.cuda(), target.cuda()

    # create new variables for hidden/cell states, otherwise
    # we'd backprop through entire training history
    hidden = tuple([c.data for c in hidden])

    # clear gradients
    rnn.zero_grad()

    # forward pass to get the output and new hidden/cell states
    out, hidden = rnn(inp, hidden)

    # output loss
    loss = criterion(out, target)

    # perform backpropagation and optimization
    loss.backward()
    nn.utils.clip_grad_norm_(rnn.parameters(), 5)
    optimizer.step()

    # return the loss over a batch and the hidden state produced by our model
    return loss.item(), hidden

```

Figure 4.10: Forward and Back Propagation in Training loop [6]

- Learning Rate: to set the learning rate of Adam Optimizer
- Vocab Size : to set the number of unique tokens in vocabulary
- Output Size : to set the desired size of output
- Embedding Dim : to set the embedding dimension, smaller than vocab_size
- Hidden Dim : to set the hidden dimension of RNN
- N Layer: to set the number of layers in the RNN


```

def train_rnn(rnn, batch_size, optimizer, criterion, n_epochs, show_every_n_batches=100):
    batch_losses = []

    rnn.train()

    print("Training for %d epoch(s)..." % n_epochs)
    for epoch_i in range(1, n_epochs + 1):

        # initialize hidden state
        hidden = rnn.init_hidden(batch_size)

        for batch_i, (inputs, labels) in enumerate(train_loader, 1):

            # make sure you iterate over completely full batches, only
            n_batches = len(train_loader.dataset)//batch_size
            if(batch_i > n_batches):
                break

            # forward, back prop
            loss, hidden = forward_back_prop(rnn, optimizer, criterion, inputs, labels, hidden)
            # record loss
            batch_losses.append(loss)

            # printing loss stats
            if batch_i % show_every_n_batches == 0:
                print('Epoch: {:>4}/{:<4} Loss: {}'.format(
                    epoch_i, n_epochs, np.average(batch_losses)))
                batch_losses = []

    # returns a trained rnn
    return rnn

```

Figure 4.11: Training Loop [6]

- Show N Batches: to set the number of batches at which neural network should print progress

```

# create model and move to gpu if available
rnn = RNN(vocab_size, output_size, embedding_dim, hidden_dim, n_layers, dropout=0.5)
if train_on_gpu:
    rnn.cuda()

# defining loss and optimization functions for training
optimizer = torch.optim.Adam(rnn.parameters(), lr=learning_rate)
criterion = nn.CrossEntropyLoss()

# training the model
trained_rnn = train_rnn(rnn, batch_size, optimizer, criterion, num_epochs, show_every_n_batches)

# saving the trained model
helper.save_model('./save/trained_rnn', trained_rnn)
print('Model Trained and Saved')

```

Figure 4.12: Adam Optimizer and Cross Entropy [6]


```
{'embedding_dim': 100, 'sequence_length': 15, 'batch_size': 100, 'hidden_dim': 512, 'n_layers': 2, 'learning_rate': 0.001, 'num_epochs': 20}
Training for 20 epoch(s)...
Epoch: 1/20    Iterations: 50000    Loss: 5.485613    Elapsed time: 39.58012938499451
Epoch: 1/20    Iterations: 100000    Loss: 4.997570    Elapsed time: 78.88181114196777
Epoch: 1/20    Iterations: 150000    Loss: 4.849201    Elapsed time: 118.13977670669556
Epoch: 1/20    Iterations: 200000    Loss: 4.733220    Elapsed time: 157.3909046649933
Epoch: 1/20    Iterations: 250000    Loss: 4.655129    Elapsed time: 196.6342270374298
Epoch: 1/20    Iterations: 300000    Loss: 4.607391    Elapsed time: 235.8708279132843
Epoch: 1/20    Iterations: 350000    Loss: 4.515004    Elapsed time: 275.098806142807
Epoch: 1/20    Iterations: 400000    Loss: 4.458964    Elapsed time: 314.35806798934937
Epoch: 1/20    Iterations: 450000    Loss: 4.424423    Elapsed time: 353.6037404537201
Epoch: 1/20    Iterations: 500000    Loss: 4.399562    Elapsed time: 392.852801322937
Epoch: 2/20    Iterations: 50000    Loss: 4.304887    Elapsed time: 449.0392520427704
Epoch: 2/20    Iterations: 100000    Loss: 4.292997    Elapsed time: 488.30241775512695
Epoch: 2/20    Iterations: 150000    Loss: 4.277061    Elapsed time: 527.5746669769287
Epoch: 2/20    Iterations: 200000    Loss: 4.233983    Elapsed time: 566.8624451160431
Epoch: 2/20    Iterations: 250000    Loss: 4.218230    Elapsed time: 606.1237514019012
Epoch: 2/20    Iterations: 300000    Loss: 4.225444    Elapsed time: 645.4023404121399
```

Figure 4.13: Training progress result [6]

4.11 Training Results

Before starting training, we are initializing: Shown in figure [4.12] and [4.13]

- The optimizer as Adam using Pytorch's library.
- loss function as Cross Entropy loss function.

4.12 Final output

The final generated script using LSTM is shown below in the figure 4.14

jerry: i mean, uh, you know, the little, uh, a couple of times. and i was in the lobby of the pool of the whole thing, i never had a little bit, you have to have it.

jerry: oh, you got a little bit of this...

george:(quietly) yeah.

kramer: hey! hey!(to elaine) what?

jerry: i can't believe that.

kramer:(laughing) what?

george: well, i don't know.

george:(to jerry and elaine) oh, come on. i got a little tired.

jerry: i think you're not getting the whole thing for you!

elaine: oh, you know what you want?

jerry: well, i don't know, i was just thinking of the way i was.

jerry: what are you doing?

helen: well, i was a little uncomfortable, i was just looking for you. and you know, they were talking, and i was just a very nervous guy, but i was in my apartment.(to jerry) you see, you know, i don't want to be a lot of water.

elaine: what are you talking about?

Figure 4.14: Generated Script using RNN [6]

5 Conclusions

5.1 Comparision of Different Architectures

According to the paper (Reference 3), the model for Tv script generation is trained in all three models based on LSTM, GRU and Bidirectional RNN. The performance of the models is further analysed to reach the conclusion that LSTM generates text in most efficient way followed by GRU and then Bidirectional RNN. Although the loss is least in Bidirectional RNN, followed by LSTM and then GRU.[3]

5.2 Final Comments

An implementation of a TV script generation employing RNN and LSTM is discussed in this work. By raising the values of several hyperparameters such as sequence length, batch size, hidden dimensions, layer count,etc. We've tried to keep the cross-entropy loss to a minimum, in the range of (9, 3) to generate structured and grammatically correct TV scripts. It is quite evident that the final script can be improved further with a better and large input data size and a sophisticated architecture and powerful resources to generate a realistic TV script. There are many applications of RNN and related architectures to work on sequential data and real-world problems like music generation, language translation, code generation, machine translation, speech recognition, video tagging, etc.

Bibliography

- [1] Chung, Junyoung, et al. *Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling*. [Arxiv, <https://arxiv.org/pdf/1412.3555v1.pdf>].December, 2014.
- [2] Dipti Pawade, et al. *Story Scrambler – Automatic Text Generation Using Word Level RNN LSTM*. [MECS, <https://www.mecs-press.org/ijitcs/ijitcs-v10-n6/IJITCS-V10-N6-5.pdf>].June, 2018.
- [3] Sandhya Mangal, et al. *LSTM vs GRU vs Bidirectional RNN for script generation*. [Arxiv, <https://arxiv.org/ftp/arxiv/papers/1908/1908.04332.pdf>].August, 2019.
- [4] Seth Weidman *Deep Learning from Scratch: Building with Python from First Principles Publisher*. O'Reilly Media, 2019.
- [5] Deep learning nano degree, Udacity
<https://classroom.udacity.com/>
- [6] Github Repository of the code
<https://github.com/binit92/INSE-6421>
- [7] Seinfeld dataset on Kaggle
<https://www.kaggle.com/thec03u5/seinfeld-chronicles#scripts.csv>
- [8] Wikipedia page of Seinfeld series
<https://en.wikipedia.org/wiki/Seinfeld>
- [9] Maths behind LSTM
<https://www.youtube.com/watch?v=iX5V1WpxxkY>
- [10] Understand LSTM Network
<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- [11] Pytorch Official Website
<https://pytorch.org/>
- [12] Numpy Official Website
<https://numpy.org/>
- [13] Pickle Official Website
<https://docs.python.org/3/library/pickle.html>

- [14] Adam Optimizer
<https://pytorch.org/docs/stable/generated/torch.optim.Adam.html#torch.optim.Adam>
- [15] Cross Entropy Loss
<https://pytorch.org/docs/stable/generated/torch.nn.CrossEntropyLoss.html>